

Trees: part 2

Petr Kurapov


Fall 2024

Agenda

- B-tree
- B+ tree
- Concurrent tree access

Motivation

- Data is organized in blocks, “low” number of pages present in the main memory – not able to store all the data in it




Latency Comparison Numbers (~2012)

L1 cache reference	0.5 ns			
Branch mispredict	5 ns			
L2 cache reference	7 ns	14x L1 cache		
Mutex lock/unlock	25 ns			
Main memory reference	100 ns	20x L2 cache, 200x L1 cache		
Compress 1K bytes with Zippy	3,000 ns	3 us		
Send 1K bytes over 1 Gbps network	10,000 ns	10 us		
Read 4K randomly from SSD*	150,000 ns	150 us	~1GB/sec SSD	
Read 1 MB sequentially from memory	250,000 ns	250 us		
Round trip within same datacenter	500,000 ns	500 us		
Read 1 MB sequentially from SSD*	1,000,000 ns	1,000 us	1 ms	~1GB/sec SSD, 4X memory
Disk seek	10,000,000 ns	10,000 us	10 ms	20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000 ns	20,000 us	20 ms	80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000 ns	150,000 us	150 ms	

Id	Col1	Col2	Col3	...
0				
1				
2				
3				

Motivation

- Data is organized in blocks, “low” number of pages present in the main memory – not able to store all the data in it
- i.e. row-store
- Sequential table scan – many accesses (# of chunks)



A chunk
fits single
block

Id	Col1	Col2	Col3	...
0				
1				
2				
3				


Motivation

- i.e. row-store
- Sequential table scan – many accesses (# of chunks)
- Use index to minimize disk accesses

Id	BlockId
0	X
1	Y


A chunk
fits single
block

Id	Col1	Col2	Col3	...
0				
1				
2				
3				



Motivation

- i.e. row-store
- Sequential table scan – many accesses (# of chunks)
- Use index to minimize disk accesses
- As index grows, fetching it becomes expensive
- Tree indexing!




A chunk fits single block

Id	BlockId
0	X
1	Y
2	Z
3	...
4	...
5	...


B(+)-tree

- Self-balancing m-way search tree
- Perfectly balanced: same height on every path from root
- Every node, except for root are at least half full
- Keys are sorted in a node




B(+)-tree

- Self-balancing m-way search tree
- Perfectly balanced: same height on every path from root
- Every node, except for root are at least half full
- Keys are sorted in a node




B(+)-tree

- Leaf nodes for B+ tree contain pointers/tuples they index
- Connected into linked list



Relation to red-black tree



B(+)-tree

B-tree

- Less space required
- Updates are more expensive for multiple thread access: rebalancing requires more latches

B+ tree

- All values are in leaf nodes
- Keys are duplicated
- Can preserve key in non-leaf node upon key removal

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent

B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent

4

B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent


5

4	8
---	---

B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent




B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent

B tree grows “upward”



B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent




B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent

B tree grows “upward”



B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion


- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”

B(+)-tree: Insertion

- Find appropriate leaf node for the new element
- If enough space – insert, preserving order
- If not, split the node in two and move middle element to parent



B tree grows “upward”


B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge

B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge


5



B(+)-tree: Deletion


- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge

5



B(+)-tree: Deletion


- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge



B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge


8



B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge


8



B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge


8



B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge


8



B(+)-tree: Deletion


- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge

8




B(+)-tree: Deletion

- Find leaf node
- If deletion doesn't make node less than half full – done
- If not, redistribute
- If unable, merge




Concurrency control

- A *latch* protects a critical section of an in-memory data structure operation
 - Read latch would allow multiple threads to read the value concurrently
 - Write latch allows exclusive access only
- Logical correctness – quacks like a duck
 - Physical correctness – the duck is quaking



Logical correctness



Physical correctness ₃₇

Concurrency control

- A *latch* protects a critical section of an in-memory data structure operation
- Read latch would allow multiple threads to read the value concurrently
- Write latch allows exclusive access only

Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#)) – compare-and-swap + OS fallback

Concurrency control

- A *latch* protects a critical section of an in-memory data structure operation
- Read latch would allow multiple threads to read the value concurrently
- Write latch allows exclusive access only

Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`


Concurrency control

- A *latch* protects a critical section of an in-memory data structure operation
- Read latch would allow multiple threads to read the value concurrently
- Write latch allows exclusive access only

Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based


Concurrency control

	Read	Write
Locked	0	0
Queued	0	0

Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based


Concurrency control



Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based

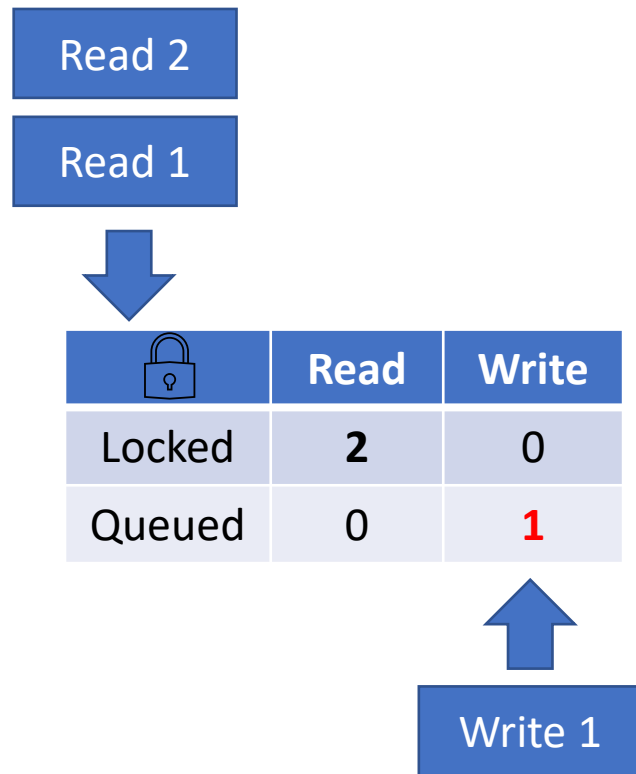
Concurrency control



Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based


Concurrency control



Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based

Concurrency control




Implementation:

- Blocking mutex
 - `std::mutex`
 - Fast user-space locking ([futex](#))
- Spinlock
 - `std::atomic`
- Read-write latch
 - Allow multiple readers storing a # of accessors
 - Policy based


B+ tree: multithreading

- Allow multiple threads to do reads and writes simultaneously
- Potential issues:
 - Concurrent writes on the same object
 - Read on split/merged data
- Thread 1: delete 6
- Thread 2: find 4




B+ tree: multithreading

- Allow multiple threads to do reads and writes simultaneously
- Potential issues:
 - Concurrent writes on the same object
 - Read on split/merged data
- Thread 1: delete 6 (stalled)
- Thread 2: find 4




B+ tree: multithreading

- Allow multiple threads to do reads and writes simultaneously
- Potential issues:
 - Concurrent writes on the same object
 - Read on split/merged data
- Thread 1: delete 6
- Thread 2: find 4




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Find 6




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Find 6




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Find 6




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Find 6




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Find 6: done




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9



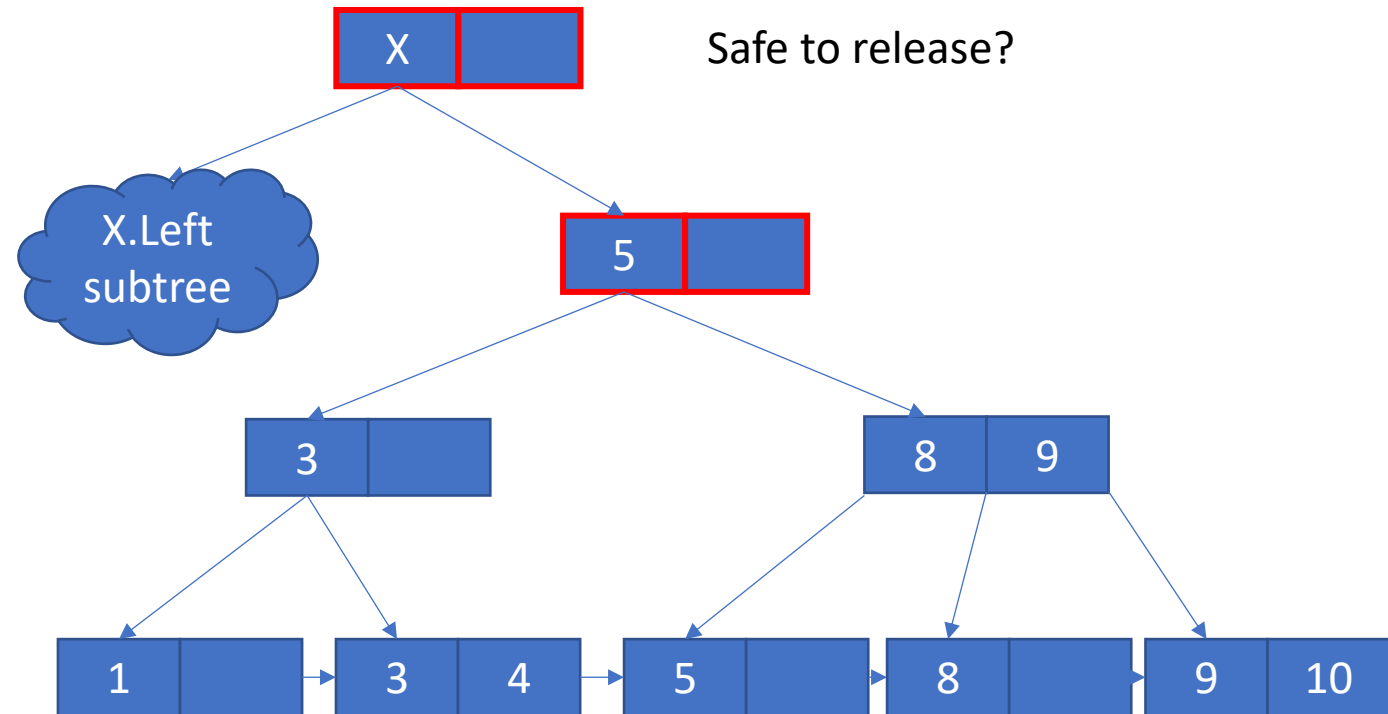
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9



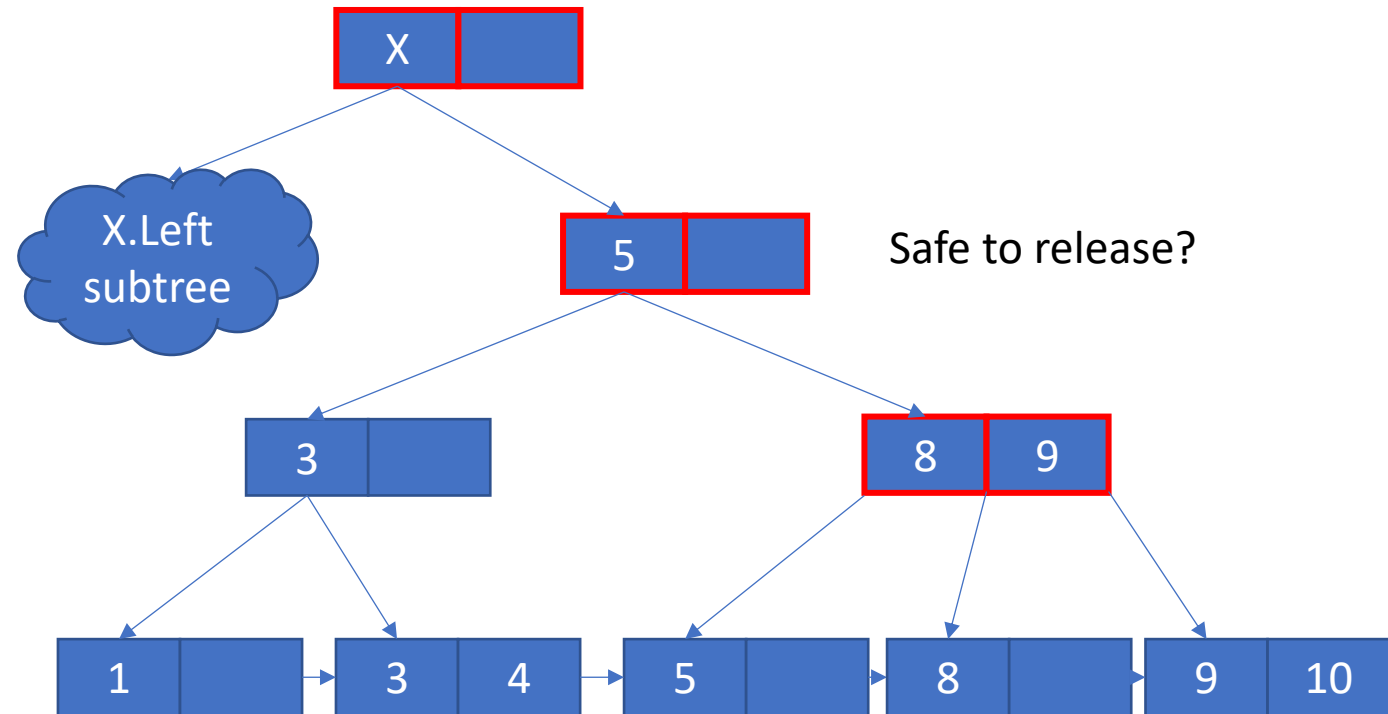
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9



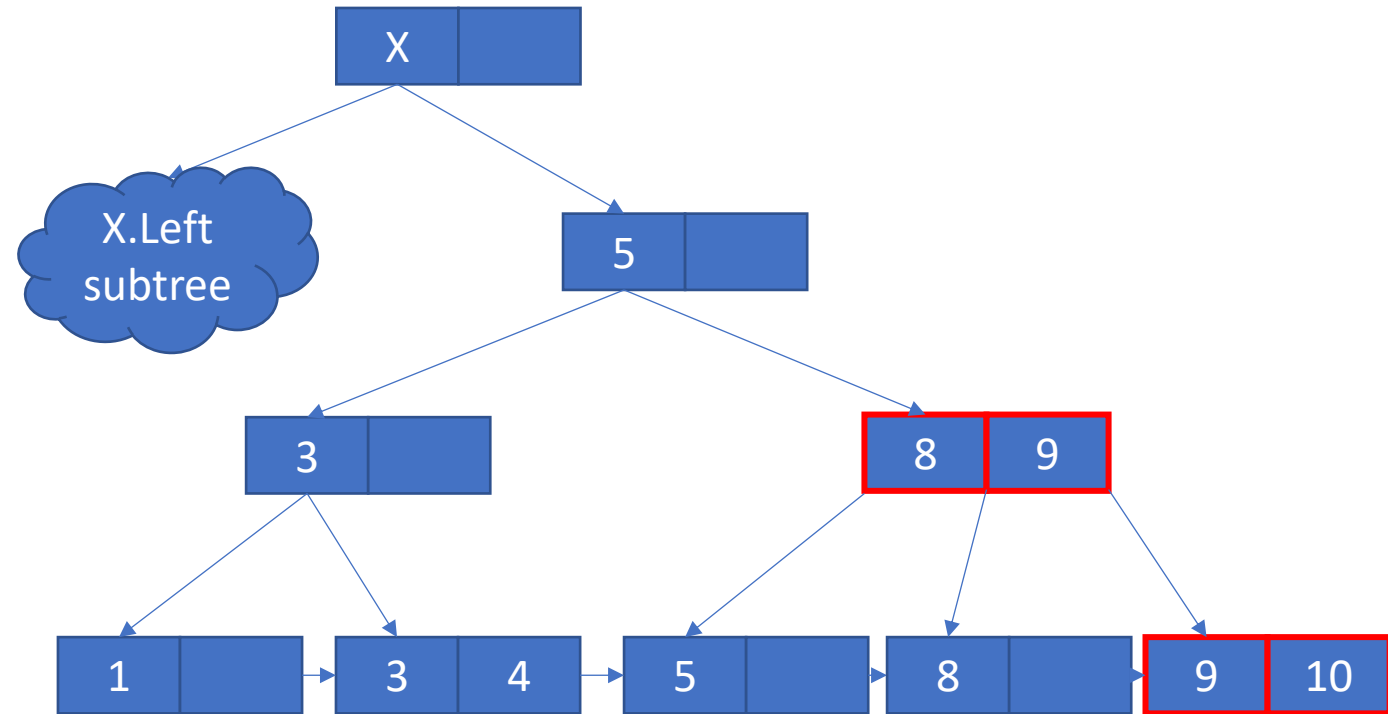
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9



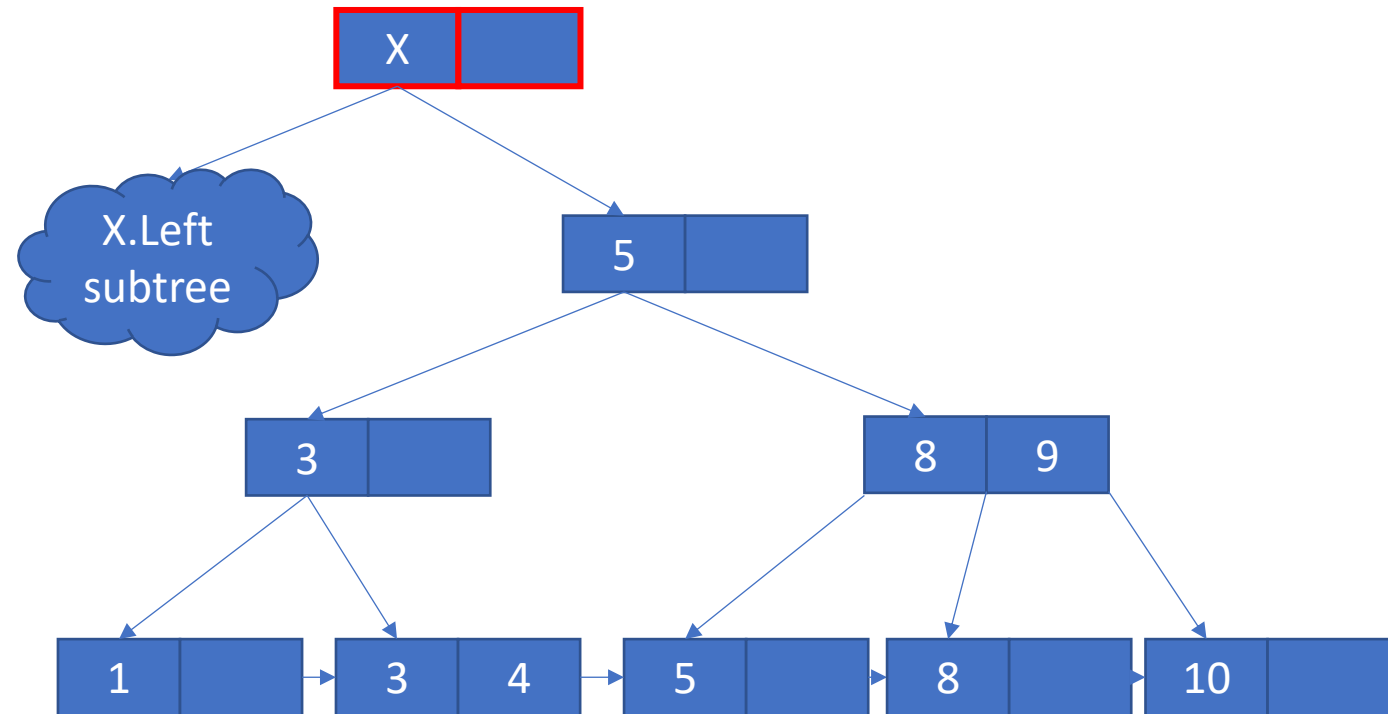
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 9



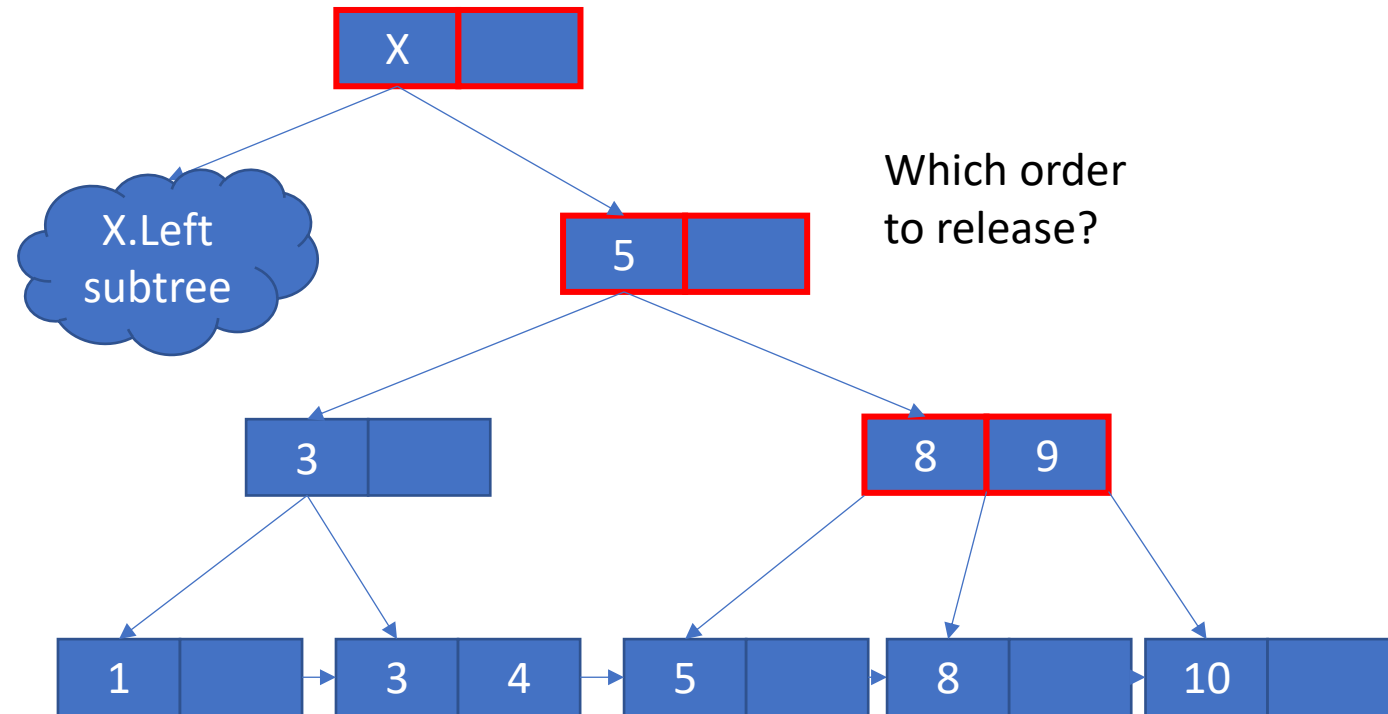
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 10



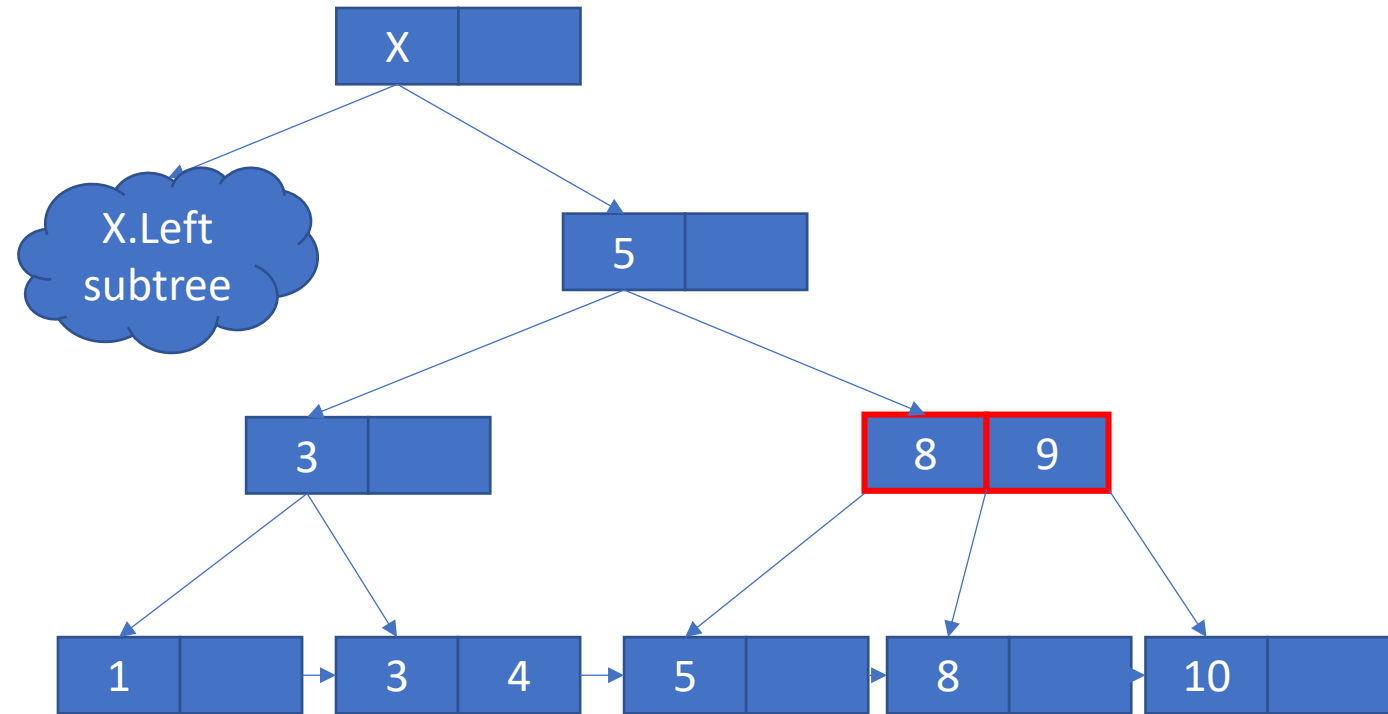
B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 10




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 10




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 10




B+ tree: multithreading

- Latch coupling
 - Acquire lock for parent
 - Acquire lock for child
 - Release parent when able to
- Delete 10: done




Anything better?

- Write-Latching root is expensive



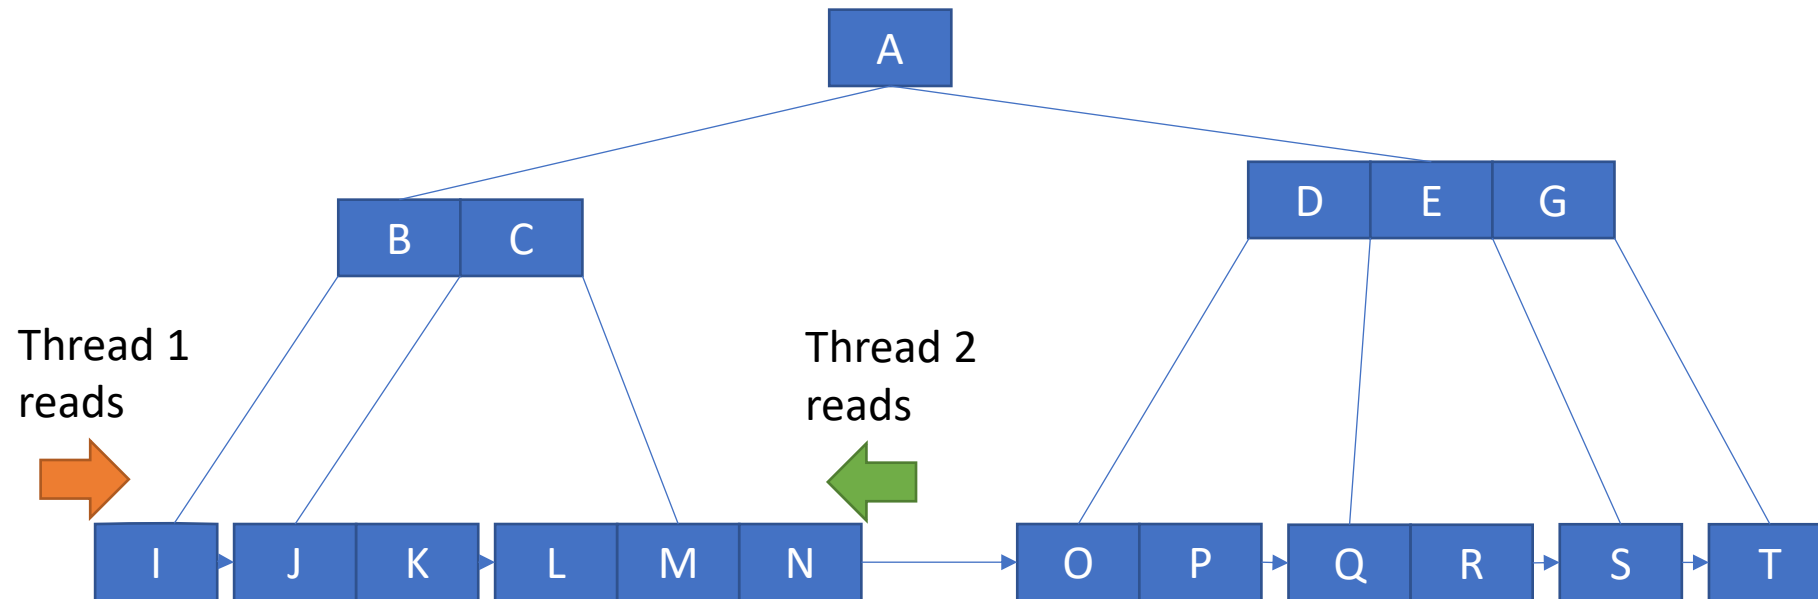
Anything better?

- Write-Latching root is expensive
- Optimistic assumption: instead of write latches use read latches
- On conflict – abort & rerun

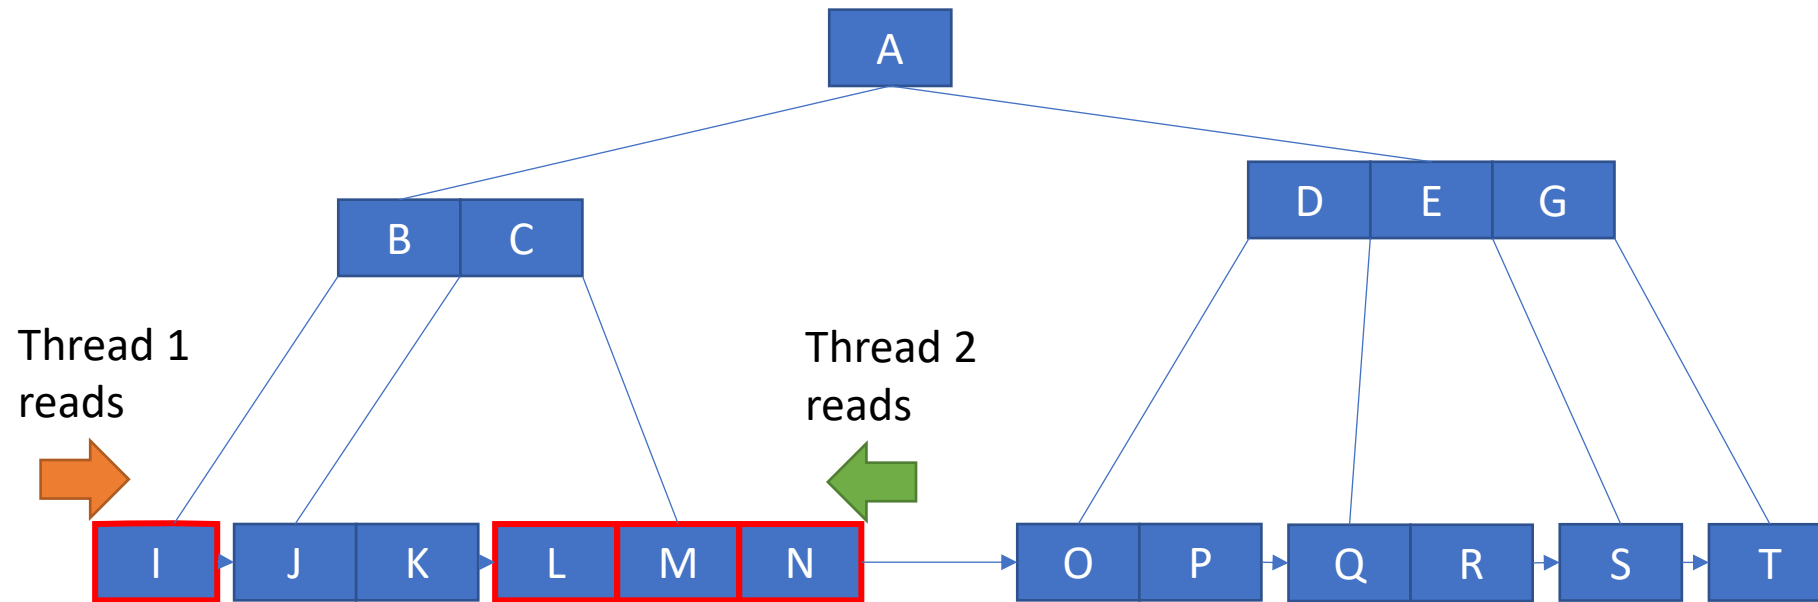


Leaves traversal

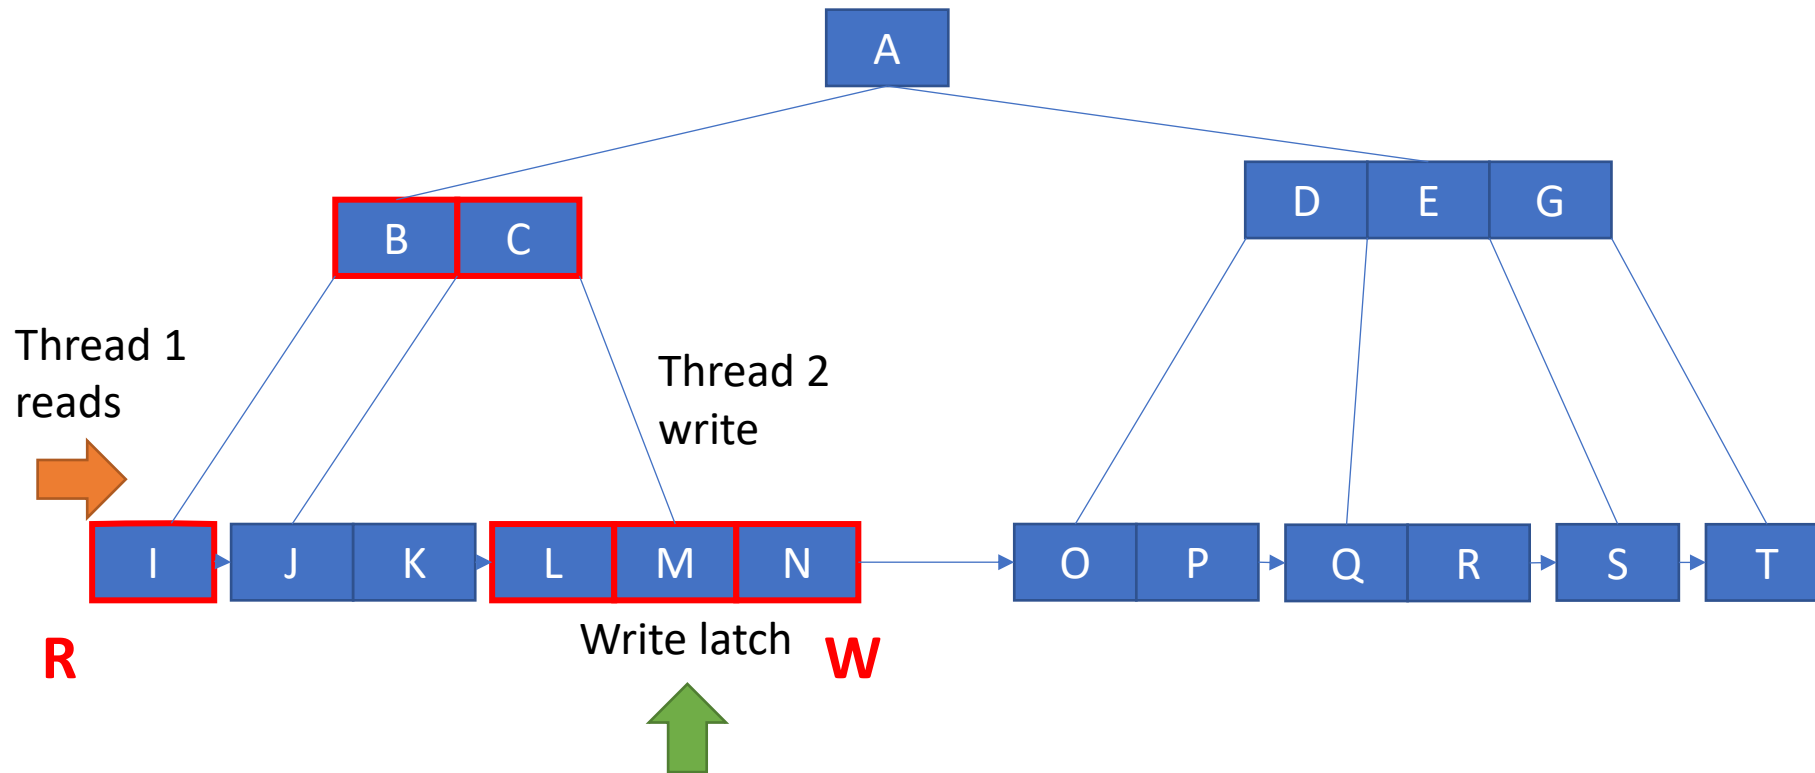
Thread 2: find all less than N



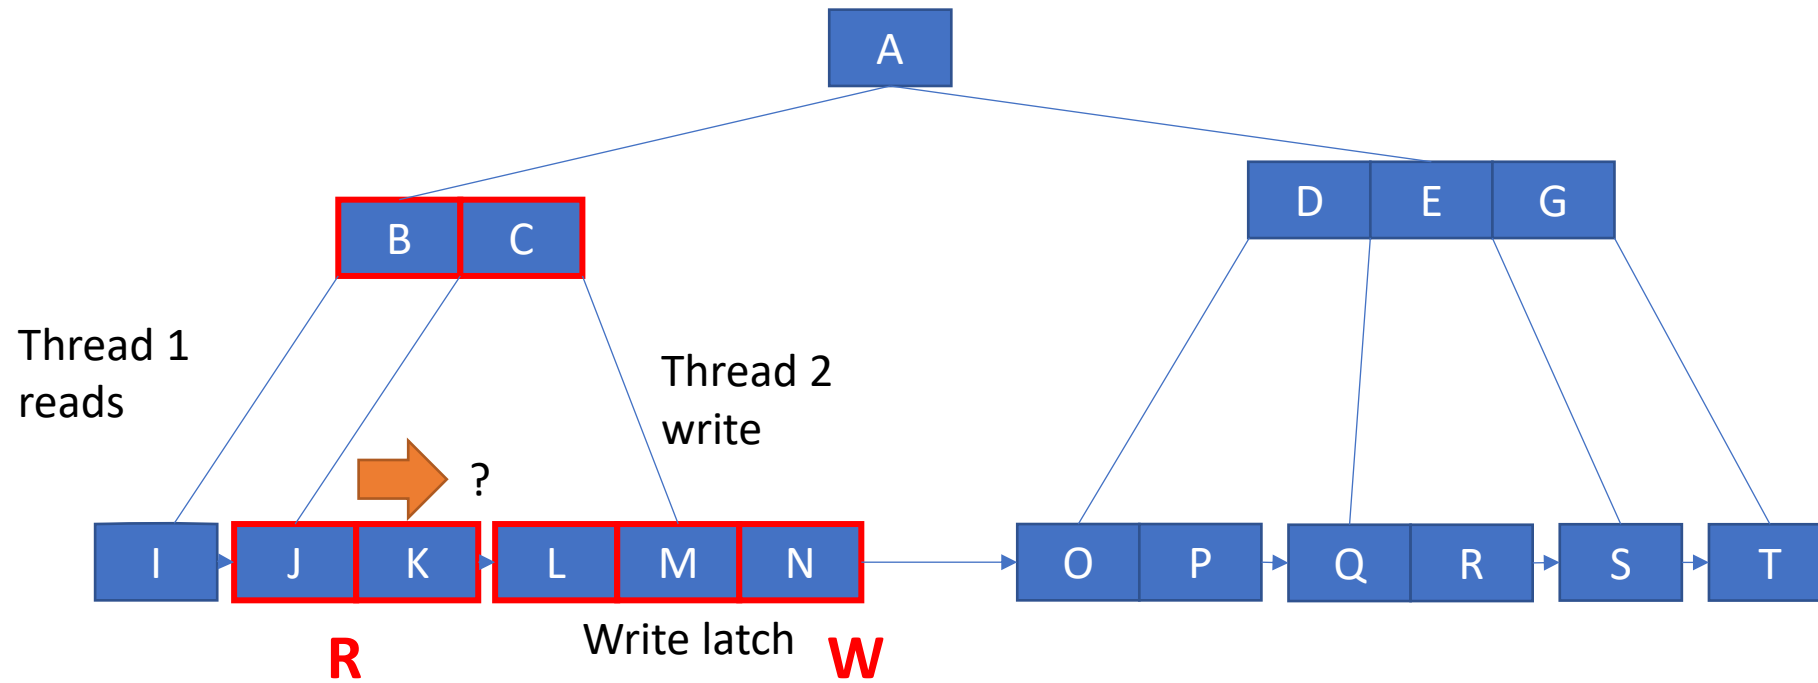
Leaves traversal



Leaves traversal

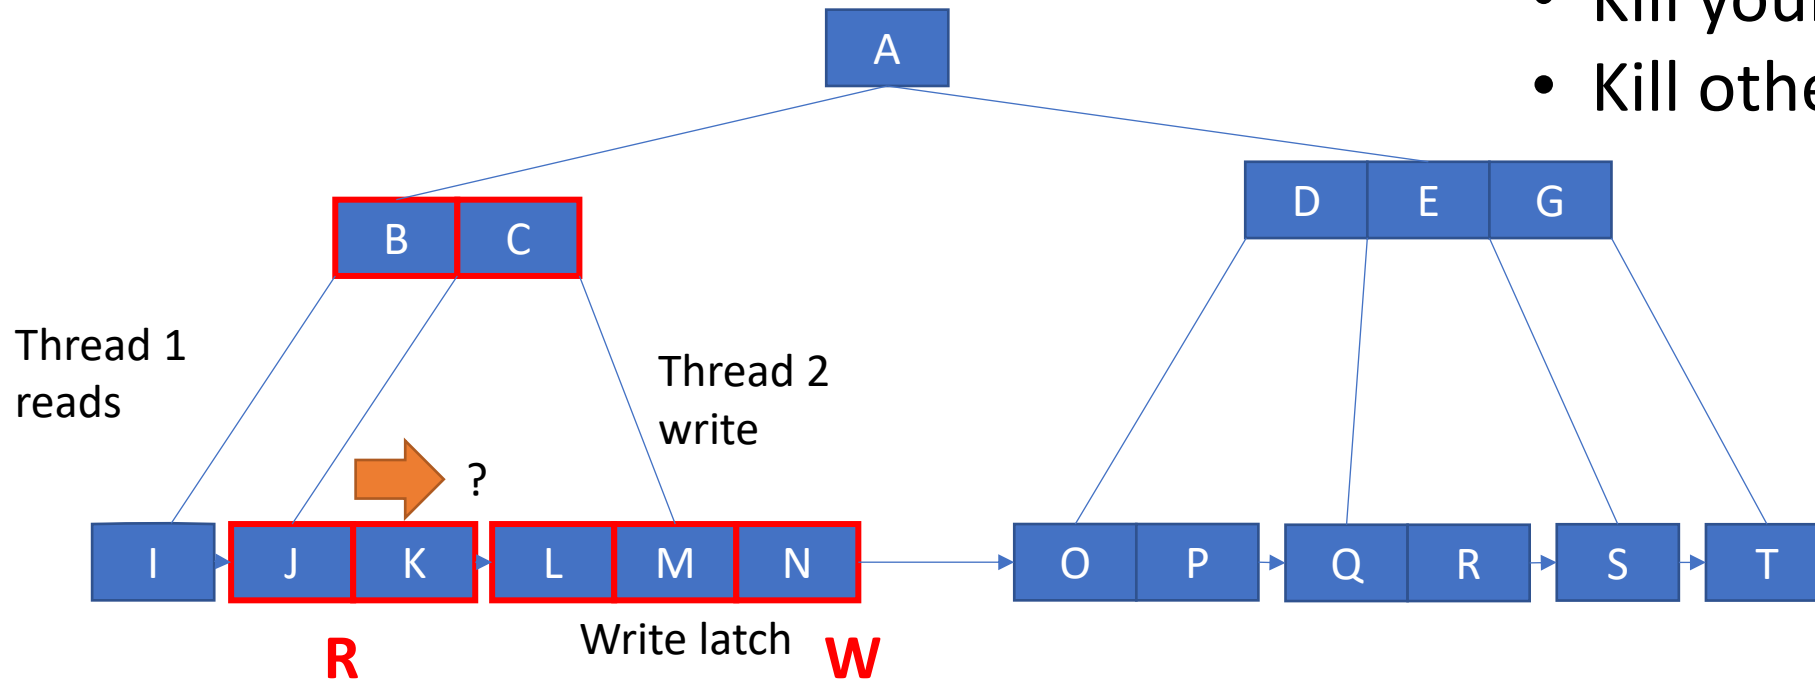


Leaves traversal



Leaves traversal

- Wait
- Kill yourself
- Kill other



Resources

- [1] Introduction to Algorithms, Thomas H. Cormen, chapters 18.
- [2] [B tree visualization](#)
- [3] [B+ tree visualization](#)
- [4] further reading: [Bw tree](#)

BACKUP